

## AI OPERATIONS

# The agent ran all night. No result.

---

Just an invoice for tokens spent retrying the same broken step.

**Run design is the fix. Not a model upgrade.**

# Your subscription is not your ceiling.

- Subscription covers access. That's all.
- API credits and reasoning tokens are separate — and they run whether the task is going anywhere or not.
- A stuck agent doesn't fail fast. It retries. It generates tokens fixing the problem it already spent tokens creating.

**It fails expensively.**

## CONTROL 1

# Set a spend cap. Before. Not after.

**Target: 60-70% of expected run cost**

If the agent burns through that without a checkpoint, something broke early.

**The buffer is information**

You want to know before you've paid full price for a broken run.

**Most platforms expose this. Most people don't touch it.**

## CONTROL 2

# Build checkpoints into the run.

## Mid-run proof of progress

A structured output after each phase, written somewhere you can inspect. If something broke at hour one, you want to know at hour one. Not hour four.

**You wouldn't manage a contractor this way. Don't manage an agent this way.**

## CONTROL 3

# Require a receipt at the end of every run.

- What task ran
- What it cost
- What it produced
- What it tried and never resolved

**A few lines in your final step. Changes your ability to improve the run over time.**



## CONTROL 4

# Give it permission to stop when the task is wrong.

**A capable agent will improvise when the spec breaks.**

It generates plausible-looking output on bad inputs. That is not what you want. A validation step early in the run — stop now if the input doesn't match the task — costs almost nothing.

**A stop condition is worth more than a bigger model.**

**THE ONE-CARD MODEL**

# One card. Four fields. Per recurring run.

- | **Max spend for this run**
- | **Check-in format and timing**
- | **What the receipt looks like at the end**
- | **When to stop and return to you**

**No new platform. No developer. Just a decision made before you hit run.**

full/REFIT

---

## Get the Agent Run Control Card worksheet.

Maps the four controls to your specific stack. Takes ten minutes to fill in once. Runs every week after that.

**Get the free worksheet**

**(Link in the first comment)**