

AI OPERATIONS

Your AI stack needs a fire drill, not brand loyalty.

If your workflow only works when one model is available,
it is not a workflow yet.

Teams are quietly building real work around:

- One account, one API key, one scheduler
- A model comparison debate nobody can settle
- A fragile fallback nobody has written down

The operating question is not which model is best.

The operating question is simpler.

What happens to the work if this model is unavailable tomorrow morning?

Most teams have not asked this out loud yet.

Pick one AI-dependent task and find out.

Run the fire drill on one task.

| What is the task?

| Which model or provider does it depend on?

| What breaks first?

| What is the fallback model?

Four questions. Write the answers down.

The last three questions hit differently.

- What is the manual fallback?
- How late can this be before it matters?
- What is the kill switch?

The kill switch is usually the most revealing part.

What a working kill switch looks like.

- | If the source cannot be read, do not summarize from memory
- | If the output has unverified numbers, do not send it
- | If a fallback degrades quality, label it before someone treats it as normal

Rules before the work breaks. Not after.

A working demo is not a dependable workflow.

Brand loyalty is fine for a hoodie.

It is a bad continuity plan.

Run the fire drill before the model goes unavailable.

What the fire drill tells you:

- Which tasks are actually fragile right now
- Where the real fallback path actually lives
- Who owns the kill switch decision

One task. Twenty minutes. No more guessing.

full/REFIT

Test one AI-dependent task for real.

The fire drill template. Seven questions. Fallback path named.
Kill switch defined.

[Download the free template](#)

[\(Link in the first comment\)](#)