

AI OPERATIONS

Your AI gets dumber the longer you use it.

Most people tune the wrong thing to fix it.

Every message you send reloads everything.

- Your rules and tool definitions
- Your full conversation history
- The model reasons across all of it, every time

As the pile grows, the answers get worse. People call it context rot.

Quality drops first. The higher bill is just what you notice.

The model is still working. It is just working harder through more noise.
The output is what changes.

Two things got worse. Only one sent you an invoice.

The fix is not "spend fewer tokens."

The fix is: stop making your model think through noise.

Two levers do most of the work. Everything else is a rounding error.

Lever one: model choice.

- | Running the wrong model for a task costs two to five times more
- | Same prompt, same output quality target, different tier
- | Stop defaulting to the most expensive one out of habit

The wrong tier is the single most common waste source.

Lever two: prompt caching.

- | Stores the stable parts of your prompt once
- | Charges a fraction of the normal rate on every message after the first
- | Up to ninety percent off the cached portion

The single biggest cost lever. It costs you nothing in quality.

Cutting this waste is not about being cheap.

- It is about giving the model less noise to reason through
- Less noise means the output stays sharper longer
- The lower bill is a side effect, not the point

The sharper answer is what you are actually after.

Before those two are right, everything else is a rounding error.

Model choice and prompt caching first. Then optimize the rest.

Run your setup through the AI Environment Scorecard to see where you stand.

full/REFIT

See where your AI setup is leaking quality.

The AI Environment Scorecard. Five minutes. Covers model tier, caching status, and prompt structure.

Run the free scorecard

(Link in the first comment)